



CMPE 685 Computer Vision

Pytorch Tutorial

Nilesh Pandey

Abu M. N. Taufique

Index

1. Numpy \leftrightarrow Pytorch
2. Basic Syntax
3. Neural Network
4. Loss and Optimization
5. Training
6. Hyperparameter : Learning Rate | Epoch
7. Hyperparameter : Suggestions

Numpy \leftrightarrow Pytorch

Pytorch is inspired by numpy operations, which can be performed on GPU.

Basic Numpy \leftrightarrow Pytorch Operation.

```
1 x = np.random.randn(5,3)
```

```
1 X_1 = torch.from_numpy(x)
```

```
1 X
```

```
tensor([[ -0.8175,  -0.4168,   0.9458],
        [  0.9310,   0.1324,  -0.1400],
        [  1.4400,  -1.1683,  -1.7375],
        [-0.1767,  -0.3250,   0.3472],
        [  0.2042,   0.8053,  -0.4884]], dtype=torch.float64)
```

```
1 Z = torch.Tensor(x)
```

```
1 Z
```

```
tensor([[ -0.8175,  -0.4168,   0.9458],
        [  0.9310,   0.1324,  -0.1400],
        [  1.4400,  -1.1683,  -1.7375],
        [-0.1767,  -0.3250,   0.3472],
        [  0.2042,   0.8053,  -0.4884]])
```

```
1 Z = Z.type(torch.float16)
```

```
1 Z
```

```
tensor([[ -0.8174,  -0.4167,   0.9458],
        [  0.9312,   0.1324,  -0.1400],
        [  1.4404,  -1.1680,  -1.7373],
        [-0.1768,  -0.3250,   0.3472],
        [  0.2041,   0.8052,  -0.4883]], dtype=torch.float16)
```

Torch.from_numpy() maintains the datatype which is important

Syntax

Basic Syntax of NN class

- Right hand code is simplest NN without any Output or any Operation.
- All code will follow the same pattern irrespective of version.

Pytorch closely follows Python, so pythonic classes can be used
And functions like general python functions.

- The 2nd example is simple neural net
With 2 layers.

Example:

nn.Conv2d
nn.BatchNorm2d
nn.ReLU

```
class VGGBlock(nn.Module):
    def __init__(self, values=***, values=***):
        super(VGGBlock, self).__init__()
        # *** Bunch of Modules and Functions ***

    def forward(self, x):
        # *** operations on the Input ***
        return out
```

```
class VGGBlock(nn.Module):

    def __init__(self, act_func=nn.ReLU(inplace=True)):
        super(VGGBlock, self).__init__()

        self.conv1 = nn.Conv2d(in_channels, middle_channels, 3, padding=1)
        self.bn1 = nn.BatchNorm2d(middle_channels)
        self.conv2 = nn.Conv2d(middle_channels, out_channels, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.act = act_func

    def forward(self, x):

        out = self.conv1(x)
        out = self.bn1(out)

        out = self.conv2(out)
        out = self.bn2(out)

        out = self.act(out)

        return out
```



VGG – 7(example)

```
import torch.nn as nn
from torchsummary import summary
```

```
class VGG(nn.Module):
    def __init__(self, num_classes=1000, init_weights=True):
        super(VGG, self).__init__()

        self.features = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d(128, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

        )

        self.avgpool = nn.AdaptiveAvgPool2d((7, 7))

        self.classifier = nn.Sequential(
            nn.Linear(256 * 7 * 7, 1024),
            nn.ReLU(True),
            nn.Dropout(),
            nn.Linear(1024, 1024),
            nn.ReLU(True),
            nn.Dropout(),
            nn.Linear(1024, num_classes),
        )

        # if init_weights:
        #     self._initialize_weights()

    def forward(self, x):
        x = self.features(x)
        x = self.avgpool(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x
```

```
1 model = VGG().cuda()
```

```
1 summary(model, (3,224,224))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,792
BatchNorm2d-2	[-1, 64, 224, 224]	128
ReLU-3	[-1, 64, 224, 224]	0
MaxPool2d-4	[-1, 64, 112, 112]	0
Conv2d-5	[-1, 128, 112, 112]	73,856
BatchNorm2d-6	[-1, 128, 112, 112]	256
ReLU-7	[-1, 128, 112, 112]	0
MaxPool2d-8	[-1, 128, 56, 56]	0
Conv2d-9	[-1, 256, 56, 56]	295,168
BatchNorm2d-10	[-1, 256, 56, 56]	512
ReLU-11	[-1, 256, 56, 56]	0
Conv2d-12	[-1, 256, 56, 56]	590,080
BatchNorm2d-13	[-1, 256, 56, 56]	512
ReLU-14	[-1, 256, 56, 56]	0
MaxPool2d-15	[-1, 256, 28, 28]	0
AdaptiveAvgPool2d-16	[-1, 256, 7, 7]	0
Linear-17	[-1, 1024]	12,846,080
ReLU-18	[-1, 1024]	0
Dropout-19	[-1, 1024]	0
Linear-20	[-1, 1024]	1,049,600
ReLU-21	[-1, 1024]	0
Dropout-22	[-1, 1024]	0
Linear-23	[-1, 1000]	1,025,000

Total params: 15,882,984

Trainable params: 15,882,984

Non-trainable params: 0

Input size (MB): 0.57

Forward/backward pass size (MB): 157.87

Params size (MB): 60.59

Loss and Optimization

```
optimizer = torch.optim.Adam(filter(lambda p: p.requires_grad, model.parameters()), lr=3e-4)
```

```
...  
filter(lambda p: p.requires_grad, model.parameters())
```

Only layers which are not frozen can be only passed to optimizer. Frozen layers are not passed to optimizer.

```
optimizer.zero()
```

optimizer accumulates loss over batch when backpropogating the loss, so it is important to empty the optimizer.

```
optimizer.step()
```

optimizer.step() performs the update on each layer, this step is performed only if we have loss present.

```
...
```

```
criterion = nn.BCEWithLogitsLoss().cuda()
```

```
...
```

Loss is the difference between the ground truth and the predicted value. Loss needs to be backpropagated, and the optimizer updates the values in each layer.

Loss and optimizer work together.|

```
...
```

Training

Training can be done as simple as the example.

The steps in Training include.

- Loop in data
- Decide GPU/CPU
- Initiate model.train()
- Initiate optimizer to zero
- Calculate Loss
- Backprop Loss
- Update the layers
- Save the model

```
model.train()

for i, (input, target) in tqdm(enumerate(train_loader), total=len(train_loader)):
    input = input.cuda()

    target = target.cuda()

    output = model(input)

    loss = criterion(output, target)

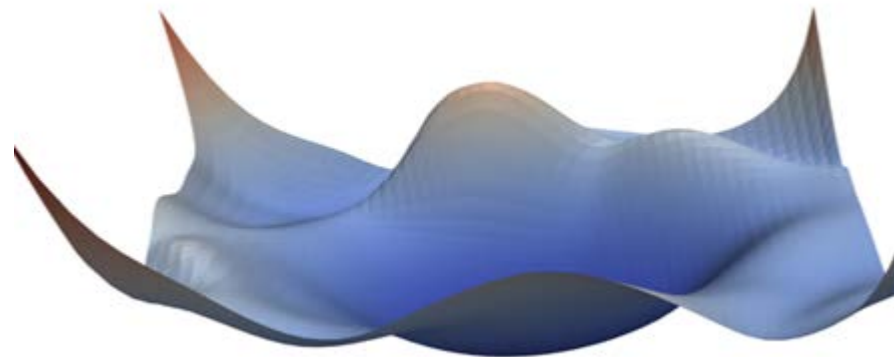
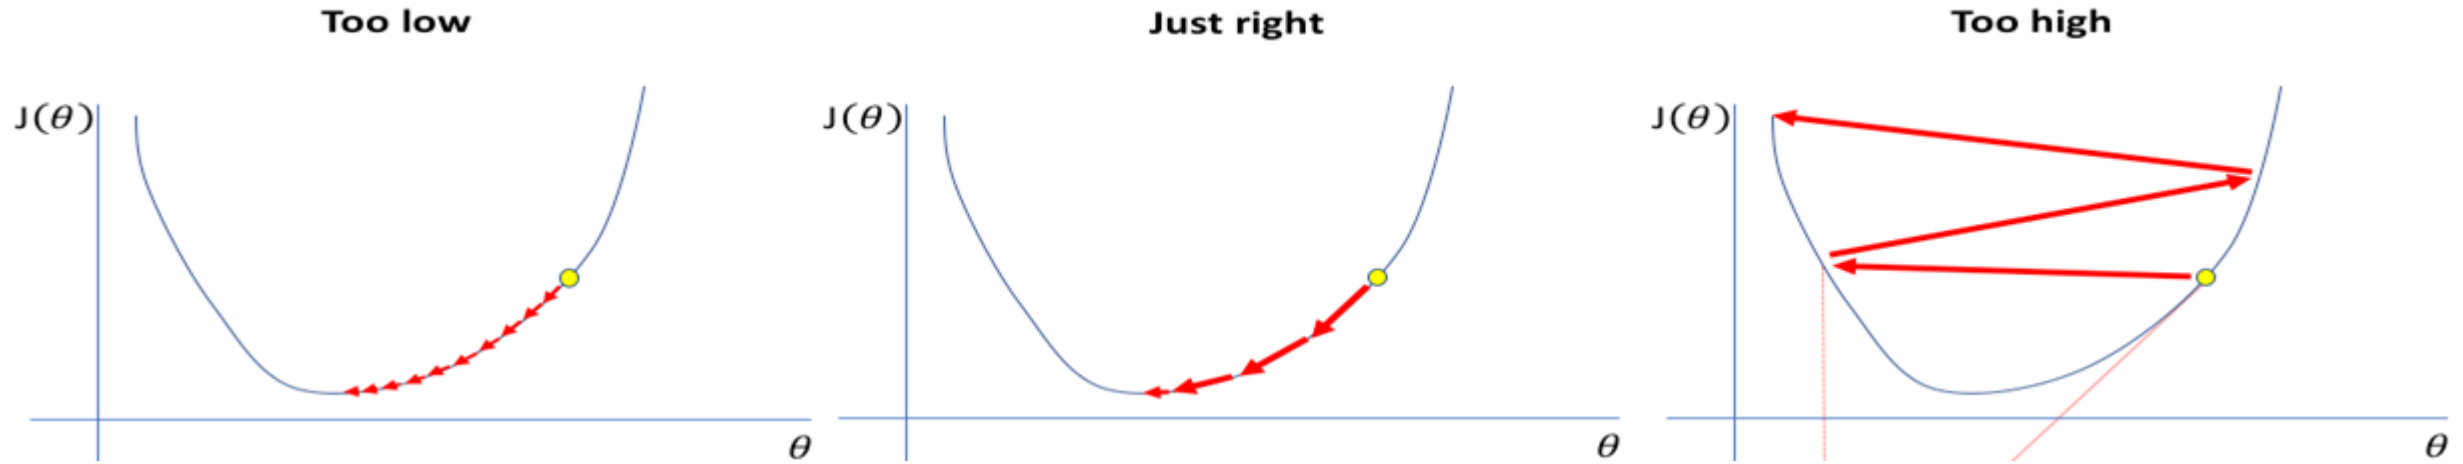
    optimizer.zero_grad()

    loss.backward()

    optimizer.step()

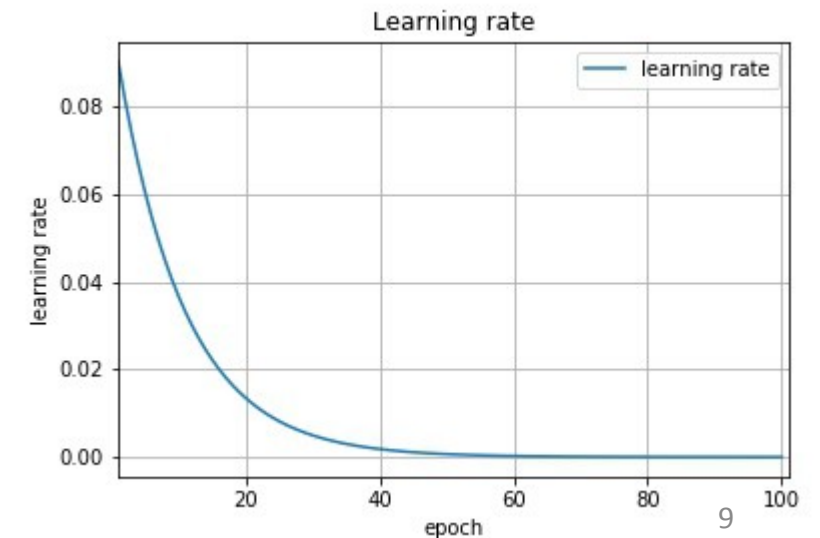
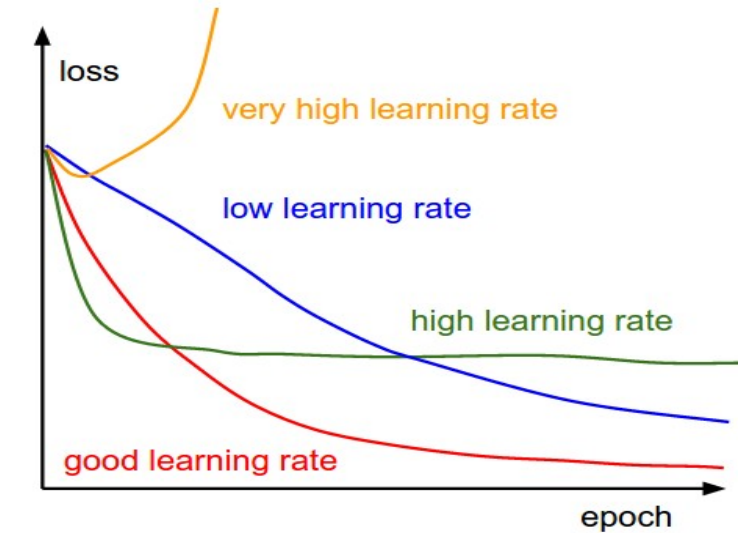
    if i%500==0:
        fig, ax = plt.subplots(1,2, figsize = (8,4))
        ax[0].imshow(target[0,0,:,:].detach().cpu(), cmap="gray")
        ax[1].imshow(output[0,0,:,:].detach().cpu(), cmap="gray")
        plt.show()
```

Hyper parameters : Learning Rate



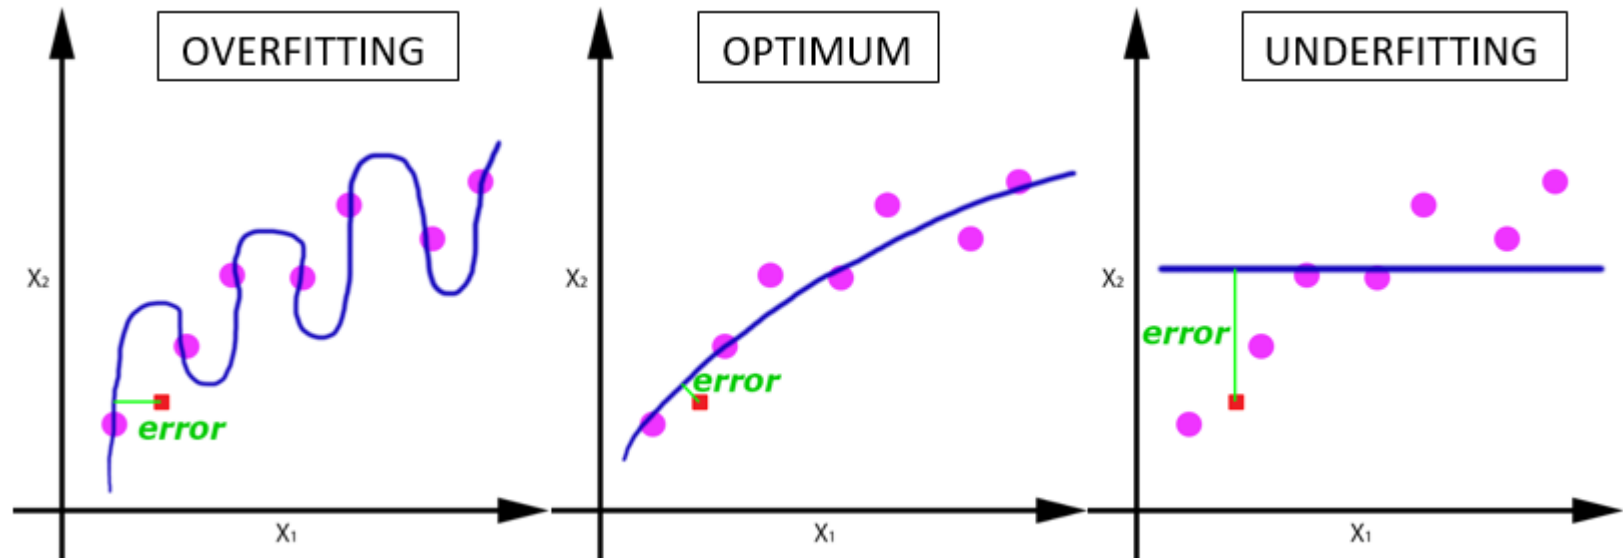
Hyper parameters : Learning Rate

- Learning rate
 1. LR based on Loss
 1. Loss \uparrow / \downarrow
 2. Convergence Should not be very fast
 2. LR Scheduling
 1. Learning Rate decay
 2. Optimizer

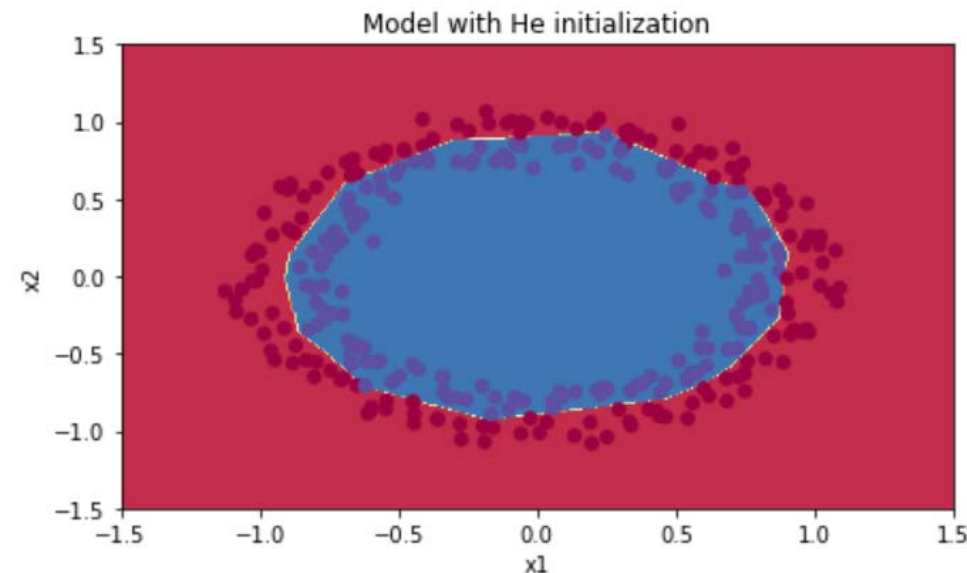
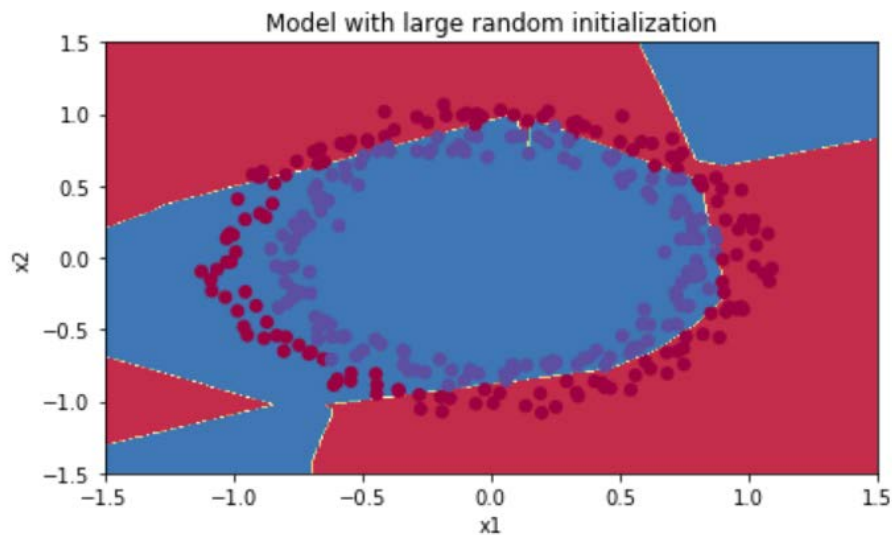


Hyper parameters : Epochs

- Higher Number of Epochs
 - First Image
- Lower Number of Epochs
 - Last Image
- Optimal Number of Epochs
 - Middle Image



Hyper parameter: Weight Initialization



Result after training network with two different weight initialization.

- Initialization is a very important hyper parameter to get better optimization.

```

'''
torch.nn.init.uniform_(tensor, a=0, b=1)
torch.nn.init.normal_(tensor, mean=0, std=1)
torch.nn.init.kaiming_normal_()
torch.nn.init.xavier_normal_()
'''

def weights_init(m):
    if isinstance(m, nn.Conv2d):
        torch.nn.init.kaiming_normal_(m.weight.data)
        torch.nn.init.kaiming_normal_(m.bias.data)

model.apply(weights_init)

```

Useful References

- <https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1>
- <https://towardsdatascience.com/activation-functions-and-its-types-which-is-better-a9a5310cc8f>
- <https://towardsdatascience.com/batch-normalization-in-neural-networks-1ac91516821c>
- <https://adeshpande3.github.io/A-Beginner%27s-Guide-To-Understanding-Convolutional-Neural-Networks-Part-2/>
- https://pytorch.org/tutorials/beginner/saving_loading_models.html
- <https://pytorch.org/docs/stable/torchvision/datasets.html>
- <https://towardsdatascience.com/understanding-learning-rates-and-how-it-improves-performance-in-deep-learning-d0d4059c1c10>
- <https://towardsdatascience.com/learning-rate-schedules-and-adaptive-learning-rate-methods-for-deep-learning-2c8f433990d1>
- <https://towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c>
- <https://towardsdatascience.com/demystifying-cross-entropy-e80e3ad54a8>
- <http://rishy.github.io/ml/2015/07/28/l1-vs-l2-loss/>